

---

# Estimator Choices and Reward Quality in GRPO: A Controlled Study on the GSM8K Benchmark

---

Yi Hou

University of Chinese Academy of Sciences  
Beijing, China  
houyi25@mailsucas.ac.cn

## Abstract

Two halves of the same question: what the reinforcement learning objective does to the reward signal, and what a weak reward signal does to the learning. I train OLMo-2-0425-1B on GSM8K with an implementation of GRPO written from scratch and run (i) prompting baselines, a standard run, a learning-rate sweep and a prompt ablation; (ii) a four-way sweep over the normalization and baseline choices that define the estimator; (iii) a deliberately off-policy regime, where each rollout batch is consumed as 32 sequential training steps, under four importance-weighting schemes; and (iv) a verifier-quality experiment in which the training reward and the test-time selector are degraded in the same, controlled way. The results I would report to a colleague are these. The prompt dominates everything else: with the zero-shot prompt the base model scores 0.08% (its outputs are degenerate continuations, not wrong answers), so almost no group carries a gradient, and the standard run only works because the 3-shot prompt reaches 18.3%. At this scale the estimator variants are indistinguishable from seed noise except that Dr. GRPO is above the standard run on both seeds. In the off-policy regime, clipping is the whole story: leaving the importance weights unclipped loses 4–6 points relative to the on-policy reference, while GRPO-style clipping and GSPO recover it entirely. Finally, the same weak verifier is far cheaper in RL than in test-time selection: a 10%-flip verifier keeps 91% and 106% of GRPO’s attainable gain across two seeds, but only 57% of best-of-32’s; a format-only verifier leaves RL with 16–30% and selection with essentially nothing; and at a 30% flip rate RL retains about seven tenths of the gain where selection retains 0.14 (measured on the internal validation split at a matched step — the test-set run of that arm did not finish). I give a mechanism for the asymmetry — the flip noise is an affine transform in expectation, and group-normalized advantages and Adam are both invariant to it — and the matched-step comparison is the piece of evidence that the prediction it makes is at least directionally right. Every number in this report is printed by a script from the run records.

## 1 Introduction

This report is an implementation and replication study of Group Relative Policy Optimization (GRPO) [11], written from scratch, that makes a 1B base model better at grade-school arithmetic. The first half is a controlled tour of the estimator: a standard run, a learning-rate sweep, a prompt ablation, four one-change variants, and an off-policy correction sweep; Section 3 reports those. Section 6 is a second experiment I designed and ran afterwards, on a question that is usually only flagged in passing: *the same verifier is used twice* — once as the training reward, once as a selector at test time — and these two uses are usually studied separately.

The framing of the second half is this. Reinforcement learning from verifiable rewards assumes the verifier is right. When it is not, the reward is a proxy, and proxy rewards are what get optimized [10, 12, 4]. For a large language model, the same proxy usually has a second job: deciding which of  $n$  sampled answers to submit. Test-time selection and RL training degrade with verifier quality, but there is no reason for them to degrade *equally*, because the mechanics differ: training averages the noisy reward over a group of samples and then takes a gradient step, while selection takes an argmax over the noisiest part of the distribution — the top of the ranking, where candidates are closest together. I test whether the asymmetry is visible at 1B scale, using a verifier whose quality I control exactly: the symbolic-equality grader used in training, with its verdicts flipped at rate  $\epsilon$ , or replaced by a check of the answer *format* alone.

### Contributions.

- A measured account of why the prompt is the first thing that matters in RLVR at small scale: zero-shot prompting leaves no learning signal at all (Section 3.1), with the mechanism visible in the reward statistics.
- A four-way estimator sweep and a four-way off-policy sweep, each with two seeds, reported as dot plots rather than means, so the reader can see where the seed spread is larger than the effect (Sections 4 and 5).
- A controlled comparison of RL and test-time selection on the same degraded verifiers and the same population of samples: 16,384 generations from 512 GSM8K questions. RL is markedly more robust; I give a mechanism (affine-invariance of the normalized advantage, plus Adam’s scale-invariance) and one falsifiable prediction (Section 6.3).
- Protocol hygiene worth stating because I got it wrong first: the same arm can differ by 8 points between the internal validation split and the test set, and the weak-verifier arms once carried their *training* reward into validation, which made an entire sweep look like something it was not (Section 2).

## 2 Setup and protocol

**Model and data.** The policy is OLMo-2-0425-1B [5], a 1B-parameter base model, trained on the GSM8K training split [3] (6,400 questions used for training, 1,024 held out for internal validation) and evaluated on the 1,319-question test split. GRPO is implemented from scratch: response log-probabilities, the group-relative advantage, the clipped surrogate objective, the aggregation, and the training step are all in my own code; the rollout engine is vLLM with an NCCL weight-transfer channel that pushes the updated policy into the server before every batch.

**The verifier.** The training reward is a symbolic GSM8K grader: it extracts the answer from a `\boxed{}` or the final line and compares it to the reference with a symbolic-equality check, and it separately reports whether the response has the required `</think> <answer>` format. All the evaluation reported here uses this same grader on the model’s outputs — never the training reward, for the reason in Section 2.

**Two decoding protocols, never mixed.** Every run is evaluated at each validation step twice: greedily (temperature 0) and by sampling (temperature 1.0). The distinction is not cosmetic. At step 0 of the standard run the two agree to within noise on some steps and differ by 21 points on others (greedy 0.423 vs. sampled 0.214); the gap narrows but never closes. The prompting baselines of Section 3.1 were taken under the sampling protocol, so every comparison between the trained models and the baselines is made in the sampled column; the greedy column is the lower-variance view and is what the sweeps in Sections 4 and 5 use. Mixing the two would turn a 5-point improvement into a 30-point one.

**Budget.** The standard run is 200 rollout steps; everything else is 80, chosen because the standard run’s accuracy is already close to its plateau by then on one seed and comparably far along on the other (Figure 1). Each configuration was run with two seeds (0 and 1); where I report one seed it is stated. Group size is 8 and the rollout batch is 256 responses (32 questions) per step, so one step collects 32 groups of 8.

Table 1: Base-model GSM8K test accuracy by prompt (1,319 questions, sampled at temperature 1.0). “Format” is the share of responses the grader accepts as well-formed; “length” is the mean response length in tokens.

| Prompt           | Accuracy      | Format | Formatted, wrong | Length |
|------------------|---------------|--------|------------------|--------|
| question_only    | 0.23%         | 17.4%  | 17.2%            | 273    |
| r1_zero (0-shot) | 0.08%         | 62.6%  | 62.5%            | 97     |
| r1_zero (3-shot) | <b>18.27%</b> | 94.2%  | 75.9%            | 108    |

**Protocol hygiene.** Three rules I adopted after violating them. (i) *Validation scores the exact verifier, whatever the training reward is.* For the first version of the weak-verifier sweep, validation reused the training reward; the format-only arms accordingly reported the format rate — near 1.0 — as their accuracy for every one of 80 steps. The curves looked like a reward-hacking story and were nothing but a wiring mistake. (ii) *Compare within one split.* The exact arm of the weak-verifier sweep scores 0.5264 on the internal validation split and 0.4443 on the test split; I use the test split everywhere and re-evaluated every checkpoint there. (iii) *Paired comparisons only.* The test split is 1,319 questions and all arms see the same questions, so a one-point difference is meaningful for a paired comparison; unpaired error bars at this  $n$  are about  $\pm 1.1$  points.

### 3 Track 1: baselines and the standard run

#### 3.1 Prompting baselines

Table 1 gives the base model’s GSM8K test accuracy under three prompts. The zero-shot results are the reason the rest of this report uses a 3-shot prompt, and they are worth reading as a failure mode rather than a number: under the zero-shot `r1_zero` prompt the model emits the answer *tags* but not answers (62.6% of responses are format-correct, 0.08% are correct), and a large fraction of the format-correct responses are degenerate continuations of the pretraining distribution — one repeats the question back, another emits HTML-ish fragments (`mode = "5" nbsp="5"`) inside a plausible-looking tag skeleton. The model has learned the shape of the output that the prompt describes without learning the task it stands for.

The arithmetic that follows from the first two rows is what fixes the experiment design: with a per-question success rate of 0.08%, the probability that a group of 8 samples is entirely wrong is 0.994, so essentially every group has zero reward variance and contributes no gradient. A 200-step run under that prompt would spend 6,400 groups to produce about 40 informative ones. The third prompt raises the per-sample rate to 18.3%, where 80% of groups carry signal. I therefore adopt the 3-shot version throughout, without changing anything else.

The third prompt in the table also has a different failure profile from the first two: `question_only` produces the longest responses (273 tokens) and the worst format rate (17.4%). It does not fail at arithmetic; it fails to enter answer mode at all, and continues the question freely (Section 3.4 returns to this arm under RL).

#### 3.2 The standard run

Figure 1 is the standard run, 200 steps, two seeds. Under the sampling protocol the two seeds go from 0.214 and 0.229 to 0.470 and 0.503 — the accuracy more than doubles — while the mean response length grows from 97 to 136 tokens (seed 0) and 99 to 126 (seed 1): the model learns to write longer chains, not to guess better. Format compliance decays slightly, from 0.987 to 0.931 and 0.981 to 0.933, so the gain is not bought by breaking the format.

Three secondary observations. (i) *The two seeds disagree about the late curve.* Seed 0 peaks at step 150 (0.553) and ends 3.3 points below its peak; seed 1 rises monotonically to the last step (0.536 at step 199, its maximum). A report that quoted only the standard run’s peak would overstate its stability. (ii) *Entropy falls fast and then stops.* Token entropy drops from 0.72 to about 0.11 within the first 120 steps and stays in the 0.10–0.12 band for the remaining 80 — low enough to be a concern, but not collapse: accuracy keeps improving in that band. (iii) *The fraction of the batch that carries a gradient is a moving target.* GRPO’s zero-variance groups are pruned, and how many survive varies

Figure 1: The standard run (200 steps, 3-shot prompt). Solid lines: greedy decoding; dashed: sampling at temperature 1.0. Both seeds are shown. The dotted line is the 3-shot prompting baseline from Table 1.





Figure 2: Learning-rate sweep (80 steps, seed 0, greedy decoding).  $3 \times 10^{-5}$  peaks at step 30 and then falls by 10 points — divergence, not slow convergence;  $3 \times 10^{-6}$  has not converged by step 80.



with the batch: 184 of 256 at the start for seed 0, 144 by its last step; 112 at the start for seed 1, 168 at the end. The nominal batch is 256, but the effective batch is meaningfully smaller and not fixed, which is worth remembering before reading a learning-rate comparison too finely.

### 3.3 Learning rate

Figure 2 sweeps the learning rate over one order of magnitude, 80 steps each.  $3 \times 10^{-5}$  is the informative one: it reaches 0.478 at step 30 and then *falls* — to 0.372 by step 80, with the sampled column falling too — which is divergence rather than the tail of a plateau.  $3 \times 10^{-6}$  is still climbing at step 80 (0.454, its last value is its maximum) and its token entropy is still 0.516, against 0.143 for the chosen rate: an under-trained run is legible in the entropy before it is legible in the accuracy. The chosen rate,  $1 \times 10^{-5}$ , ends at 0.526 and is the on-policy reference used by the rest of this section.

### 3.4 The prompt ablation under RL

Figure 3 runs the standard configuration again with the bare-question prompt, whose per-sample base rate is 0.23%. Two things happen. First, the arm is extremely seed-dependent: seed 0 climbs from 0.000 to 0.089 — still rising at step 80 — while seed 1 sits at 0.013. At the start of training each step contains, in expectation, about 0.6 correct responses out of 256, so which few responses happen to be

Figure 3: `question_only` under GRPO (80 steps, both seeds, greedy decoding). The same configuration that more than doubles accuracy with a 3-shot prompt barely moves here, and the two seeds disagree by a factor of seven.





Figure 4: Four one-change variants of the standard estimator, 80 steps, both seeds, final greedy accuracy. Each point is one seed; the bar is the two-seed mean. The dashed line is the standard run’s own 80-step value (0.526, seed 0). Variant definitions in the text.




correct dominates the gradient; the two seeds are effectively sampling different problems. Second, seed 0’s trajectory is a real result in miniature: RL against *no examples* recovers about half of what the 3-shot prompt provides (0.089 vs. 0.183), by learning the answer format (the boxed-answer rate rises from 0.23 to 0.69) and then some arithmetic, from a starting point where the format does not exist. With two seeds this arm cannot support a claim about how far it would go, and the honest summary is the one the figure shows: sometimes it works, and when it works it is slow.

## 4 Track 1: estimator variants

The standard estimator normalizes each sequence’s loss by its length and divides the group’s advantages by their standard deviation. Figure 4 changes the normalization choices one at a time: `grpo_constant` replaces the length normalization with a constant, `dr_grpo` additionally drops the standard-deviation division [9], `rft` additionally removes the group baseline so that only above-average samples are reinforced, and `maxrl` divides by the group mean instead of the standard deviation. (Each choice reweights the difficulty range of the groups: dividing by  $\hat{\sigma}$  re-weights each group by  $(\eta(1 - \eta))^{-1/2}$ , which up-weights the extremes — groups that are nearly all right or nearly all wrong — relative to the middle.)

Figure 5: Off-policy corrections (80 steps, both seeds, final greedy accuracy). Each run consumes a rollout batch as 32 sequential training steps, so the last step trains on data 32 steps stale. Only the correction scheme differs.



The result is a negative one, and the figure is drawn to show why: the four variants span 0.046 in two-seed means, while the two seeds of a *single* variant differ by 0.018–0.030. Only Dr. GRPO is above the standard run on both seeds (0.529/0.544 against 0.526), and by a margin smaller than the seed spread of any other variant. With two seeds I can say that the normalization and baseline choices do not matter much at this scale and horizon — not that they do not matter, because 80 steps at 1B parameters and 18% accuracy is a regime where every variant is still moving. The one consistent ordering, Dr. GRPO ahead, is also the one the derivations predict (it is the variant that removes both the length and the variance reweighting).

## 5 Track 1: off-policy corrections

The last track-1 experiment makes the training deliberately off-policy. The rollout batch stays 256, but the *training* batch is 8 with no gradient accumulation, so the batch is consumed as 32 consecutive optimizer steps; by the end, the data is 32 policy updates stale. Four runs change only how the staleness is corrected: NAIVE ignores it (the objective is computed as if the data were on-policy), NOCLIP multiplies each token’s term by the importance ratio  $\pi/\pi_{\text{old}}$  without clipping, CLIP is the standard GRPO clipped surrogate with clip range 0.2, and GSPO uses the sequence-level ratio with a much tighter range [15].

Figure 5 shows the split clearly. The two runs that use a clipped ratio sit at the on-policy reference (0.538/0.521 and 0.523/0.531 against 0.526 — no measurable cost to  $32\times$  staleness), while the two that do not lose 3.5–6.4 points (0.481/0.491 and 0.480/0.463). Within that second pair, the run with *unclipped* weights is the worse of the two on one seed: multiplying by an unbounded ratio amplifies the few tokens whose ratio has drifted far from 1, which is a larger error than ignoring the staleness altogether. One seed is not evidence, but the direction — clipping is load-bearing, not decoration — matches the mechanism, and the contrast between the clipped and unclipped columns is much larger than the seed spread.

## 6 Track 2: what a weak verifier costs

**Design.** The verifier has two uses: it provides the training reward, and it picks the answer when several are sampled. I degrade one verifier in one way and measure both uses against the same ground truth. The degradation is applied to the symbolic grader above: *flip* its verdict with probability  $\epsilon$  (“noisy”), or keep only the format half of the reward (“format-only”). The sample population is fixed: 512 GSM8K questions, 32 samples each, generated once by the base model with the 3-shot prompt (16,384 responses). Selection re-ranks exactly that population; training uses the same verifiers as the reward and is evaluated, always, with the true grader on the full 1,319-question test split.

Figure 6: Test-time selection over the fixed sample population (512 questions, 32 samples each). Best-of- $n$  re-ranks the same samples under verifiers of decreasing quality; “oracle” is the coverage ceiling (pass@ $n$ , the unbiased estimator of [2]). The horizontal line is the no-selection accuracy ( $n=1$ ).





Figure 7: Training with a weak verifier (80 steps, both seeds), evaluated with the true grader on the 1,319-question test split. Each point is one seed, the bar is the mean; the dotted line is the untrained base model under the same protocol. The 30%-flip arm was run but its test-set re-evaluation did not complete, so it is absent here and reported as unfinished in Section 8.




## 6.1 Selection: the ceiling, and what degradation costs

Figure 6 is the panorama. With a perfect verifier, best-of-32 reaches 0.772, which is exactly the coverage ceiling (pass@32 = 0.7715): *selection can realize everything the samples contain*. Majority voting, which needs no verifier, gets 0.455 — better than the base rate of 0.209 but far below the ceiling. In between, the verifier’s reliability sets the height: with 10% of verdicts flipped, best-of-32 reaches 0.531; at 30%, 0.287; at 50%, 0.211 — statistically indistinguishable from not selecting at all. And with a format-only verifier, 0.213: a verifier that measures something real but *wrong* (format correlates with correctness only weakly) is as useless as a coin flip at selecting, which is the cleanest illustration of the distinction between a verifier being informative about the response and informative about the answer.

## 6.2 Training: the same verifiers as the reward

Figure 7 runs the standard 80-step configuration three ways, changing only the training reward, and evaluates every arm with the true grader on the test split. The exact arm reaches 0.444, with the

Figure 8: Retained fraction of the attainable gain as the verifier degrades, for the two uses of the same verifier. Selection: best-of-32 minus the base rate, over the coverage ceiling minus the base rate, on the 512-question population of Figure 6. RL: the trained arm minus the *same-protocol* prompting baseline, over the exact arm minus that baseline, on the test split; each dot is one seed. The RL series stops at  $\epsilon = 0.1$  because the 30%-flip arm has no test-set evaluation — that case is compared on the internal validation split in Section 6.2. The dashed line is the closed form for the expected-gradient scaling of RL under flip noise,  $(1 - 2\epsilon)$ .



3-shot prompting baseline at 0.183, so the attainable gain on this axis is 0.261. Against that yardstick: the 10%-flip arm keeps 91% (seed 0, 0.421) and 106% (seed 1, 0.459) — i.e. nothing measurable is lost to a 10% flip rate — while the format-only arm keeps 30% and 16% (0.260, 0.225). The format-only arm is worth a sentence of detail because it is the pure case of proxy optimization: it reaches a format rate of 0.97, so its reward is nearly maximal, while its accuracy sits a third of the way from the base model to the exact arm. The model optimizes exactly what the reward measures; a verifier that measures the wrong thing does not slow the model down, it points it elsewhere.

A third noise level,  $\epsilon = 0.3$ , was run but never evaluated on the test set: seed 0 completed its 80 steps and seed 1 reached step 67 before the shared cluster became unusable, and the re-evaluation that would put the arm in the table above never ran. It can still be compared one protocol over. On the internal validation split — 1,024 questions, scored by the exact grader — every run has a step-60 point, before their lengths diverge: the exact arm is at 0.5166 and the 10%-flip arm at 0.5156, while the 30%-flip arm’s two seeds sit at 0.4805 and 0.5010. Against the arms’ shared starting point (0.4229) and the exact arm’s gain, that is 99% retained at  $\epsilon = 0.1$  and 61% / 83% at  $\epsilon = 0.3$ . The cost is small at 0.1 and real, though modest, at 0.3 — weak support for the prediction in Section 6.3, with the caveats that the two 30%-flip seeds differ by two points and that only one  $\epsilon = 0.1$  seed is usable on this split (the other was scored with the wrong grader before I fixed it; Section 2).

### 6.3 The asymmetry, and why it should exist

Figure 8 puts the two halves on one axis. At  $\epsilon = 0.1$  the same noise costs selection 43% of its attainable gain and RL nothing measurable (the two seeds land at 0.91 and 1.06 of the exact arm). The gap widens with noise: selection retains 14% at  $\epsilon = 0.3$  and essentially zero at 0.5, while the RL series stops at  $\epsilon = 0.1$  because the  $\epsilon = 0.3$  arm’s evaluation did not complete.

Two mechanisms plausibly explain the asymmetry, and they point the same way.

**1. RL averages; selection takes an argmax.** A group of 8 rollouts is scored, and the update uses the group’s mean reward as the baseline. Flipping one verdict moves that group’s contribution by roughly  $1/8$  of a reward unit, so the noise enters the gradient already averaged over 8 samples, and then over the 32 groups of a step, and then over 80 steps: an unbiased error of this kind shrinks like  $1/\sqrt{\text{number of samples}}$  in the quantities that matter, and as a *shift* of the reward level it disappears entirely (below). Selection has no such averaging: it keeps the single highest-scored candidate, and the candidates near the top of the ranking are the ones whose scores are closest together. At  $n = 32$ , the gap between the best and second-best true scores is small, so a flip probability of  $\epsilon$  per candidate

translates into a much larger probability of choosing the wrong winner — an order statistic, and one whose sensitivity grows with  $n$ .

**2. The expected effect of symmetric flip noise is an affine transform, and GRPO is immune to affine transforms.** If the true correctness indicator is  $\eta \in \{0, 1\}$  and the verdict is flipped with probability  $\epsilon$ , then

$$\mathbb{E}[\tilde{r} \mid \eta] = (1 - \epsilon)\eta + \epsilon(1 - \eta) = (1 - 2\epsilon)\eta + \epsilon, \quad (1)$$

so the expected reward is an affine function of the true one, and  $\nabla_{\theta} \mathbb{E}[\tilde{r}] = (1 - 2\epsilon)\nabla_{\theta} \eta$ : the noise scales the gradient and shifts the reward level, and it does not move the optimum. Now note what GRPO does with a reward vector: the advantage is  $A_i = (r_i - \bar{r})/\hat{\sigma}$ , which is invariant to  $r \mapsto ar + b$  for any  $a > 0$ . The  $+\epsilon$  shift cancels in the numerator and the  $(1 - 2\epsilon)$  scale cancels in both numerator and denominator. Adam compounds this: its update is invariant to a global rescaling of the gradient as well (the first and second moments scale together). So the first-order effect of the closed form is neutralized *exactly* by the estimator’s normalization — what remains is the per-sample variance the flips inject, which is second order. My observation is consistent with this: the 10%-flip arm does not land at the 80% that the closed form might suggest it should lose (it loses nothing measurable), because the closed form is about the expected gradient, and the estimator normalizes exactly that away.

The second mechanism also makes a prediction, which is why the 30%-flip arm was run: if the residual cost is the injected variance rather than the scaling, then it should grow faster than  $(1 - 2\epsilon)$  once the flip rate is high enough for variance to dominate — an arm at  $\epsilon = 0.3$  should retain markedly less than the 91–106% of the 10% arm. The matched-step validation comparison in Section 6.2 is the evidence I have on this: about 72% retained across the two 30%-flip seeds, against 99% at  $\epsilon = 0.1$  on that split. The direction matches the prediction; the test-set version of the comparison is the experiment the cluster did not let me finish.

**Caveats on this axis.** (i) The exact arm has a single seed, so the denominators in Figure 8’s RL numbers carry no seed variance; what is robust is that both flip-arm seeds land at or above the exact arm. (ii) The TTS curve and the RL points are computed on different sets (512 questions with 32 samples vs. 1,319 questions with 1 sample), so their absolute levels should not be compared — only their retained fractions. (iii) The verifier here is a program with a controlled failure mode, not a learned reward model; the numbers are properties of this protocol, and the transferable statement is the ordering between the two uses, not the specific retained fractions. (iv) Format-only degrades the two uses differently but both severely, so the  $\epsilon = 0.1$  result is the sensitive one and it is a single noise level; the 30% arm is the test.

## 7 Related work

**RL with verifiable rewards.** GRPO [11] introduced the group-relative advantage used throughout this report; the R1 recipe [6] and open replications popularized the setting. My estimator sweep tests a subset of the objective variants that have been proposed since: dropping the standard-deviation division [9], constant loss normalization and decoupled clipping [14], sequence-level importance ratios [15], and rejection-sampling-style updates [7]. The honest reading of Section 4 is that at 1B parameters and 80 steps these choices are within seed noise of each other, with Dr. GRPO marginally ahead on both seeds.

**Verifiers and their failure.** Training verifiers to re-rank samples is the original use case for verifiers in this problem set [3]; process-level supervision is more effective but costlier [8, 13]. Reward-model over-optimization — a proxy reward going up while the true objective falls — is well documented at larger scale [4] and is the general phenomenon that the format-only arm of Figure 7 reproduces in miniature. The asymmetry I measure is a comparison between two *uses* of one degraded verifier; I have not seen that comparison isolated elsewhere, but the components — verifier quality, selection, and proxy optimization — are each well studied, and the result should be read as a controlled instance of the general warning rather than a new phenomenon.

**Test-time selection.** Best-of- $n$  with a verifier, and repeated sampling with a majority vote, are standard test-time-compute strategies; the coverage ceiling measured by  $\text{pass}@n$  uses the unbiased estimator of [2], and the observation that repeated sampling with selection converts coverage into

accuracy is the “large language monkeys” effect [1]. The ceiling in Figure 6 matches the estimator exactly at  $n = 32$ , which is the sanity check that makes the degraded-verifier curve interpretable.

## 8 Limitations

Beyond the caveats stated in place: the model is 1B parameters and the task is grade-school arithmetic, so nothing here is a scaling result; every configuration has two seeds, and I have been explicit that differences below a few points are not resolved by them; the runs are 80 or 200 steps, and both the standard run and the  $3 \times 10^{-6}$  learning rate are still improving at their last step, so these are snapshots rather than converged values; and the experiment’s verifier is a program with an injected failure mode, which makes the mechanism reproducible but leaves open how much of it carries over to a learned reward model with correlated errors. One experiment was run and not finished on its intended protocol: the  $\epsilon = 0.3$  flip arm of Section 6.2. Its two runs stopped at 80 and 67 steps when the shared cluster failed, and the test-set re-evaluation that makes an arm comparable to that table never ran. It is therefore absent from Figure 7 and from the test-set numbers, and appears only as the matched-step internal-validation comparison in Section 6.2 — same split, same grader, same step, and reported as weak support rather than folded into a table it does not belong in. One experiment I designed and did not run: a *self-verification* arm, where the policy’s own verdict supplies the reward. A 1B model that produces well-formed tags only 17–62% of the time is not a credible judge, the verdict prompt and the unparseable-answer convention would be free parameters of the result, and the outcome — verdicts with little discriminative power — is already implied by the format-only arm. I would rather state that reasoning than report the number it predicts.

## 9 Conclusion

At 1B scale, on GSM8K, with an estimator implemented from scratch: the prompt is the first-order decision (0.08% versus 18.27% base accuracy decides whether any gradient exists at all); the reward normalization and baseline variants are second-order (within seed noise over 80 steps, with Dr. GRPO weakly ahead); clipping in the importance-weighted surrogate is first-order when training is off-policy (4–6 points, versus none for the clipped forms); and the quality of the verifier — the thing RLVR actually optimizes — costs the two uses of that verifier very differently, with training retaining nearly all of the attainable gain under noise that destroys most of selection’s. The last result has a clean mechanism in the affine-invariance of the normalized advantage, and a running experiment that can falsify it.

## References

- [1] Bradley Brown, Jordan Juravsky, Ryan Ehrlich, Ronald Clark, Quoc V. Le, Christopher Ré, and Azalia Mirhoseini. Large language monkeys: Scaling inference compute with repeated sampling. *arXiv:2407.21787*, 2024.
- [2] Mark Chen, Jerry Tworek, Heewoo Jun, et al. Evaluating large language models trained on code. *arXiv:2107.03374*, 2021.
- [3] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, et al. Training verifiers to solve math word problems. *arXiv:2110.14168*, 2021.
- [4] Leo Gao, John Schulman, and Jacob Hilton. Scaling laws for reward model overoptimization. *ICML*, 2023.
- [5] Dirk Groeneveld, Iz Beltagy, Pete Walsh, et al. OLMo: Accelerating the science of language models. *ACL*, 2024.
- [6] Daya Guo, Dejian Yang, Haowei Zhang, et al. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv:2501.12948*, 2025.
- [7] Nathan Lambert, Jacob Morrison, Valentina Pyatkin, et al. Tülu 3: Pushing frontiers in open language model post-training. *arXiv:2411.15124*, 2024.

- [8] Hunter Lightman, Vineet Kosaraju, Yura Burda, et al. Let’s verify step by step. *arXiv:2305.20050*, 2023.
- [9] Zichen Liu, Changyu Chen, Wenjun Li, et al. Understanding R1-Zero-like training: A critical perspective. *arXiv:2503.20783*, 2025.
- [10] Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. Concrete problems in AI safety. *arXiv:1606.06565*, 2016.
- [11] Zhihong Shao, Peiyi Wang, Qihao Zhu, et al. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv:2402.03300*, 2024.
- [12] Joar Skalse, Nikolaus H. R. Howe, Dmitrii Krasheninnikov, and David Krueger. Defining and characterizing reward hacking. *NeurIPS*, 2022.
- [13] Jonathan Uesato, Nate Kushman, Ramana Kumar, et al. Solving math word problems with process- and outcome-based feedback. *arXiv:2211.14275*, 2022.
- [14] Qiying Yu, Zheng Zhang, Ruofei Zhu, et al. DAPO: An open-source LLM reinforcement learning system at scale. *arXiv:2503.14476*, 2025.
- [15] Chujie Zheng, Shixuan Liu, Mingze Li, et al. Group sequence policy optimization. *arXiv:2507.18071*, 2025.